



# Leads Case Study Presentation

---

BY -

DHAWAL ARORA

HEMANSHU GHADIGAONKAR


SHRINIDHI RAMESH

*Your best quote that reflects your approach... “It’s one small step for man, one giant leap for mankind.”*

- NEIL ARMSTRONG

# Introduction

---



In this case study, I examined the poor lead conversion rate of X Education, a company that markets online courses through websites and search engines. My goal was to create a logistic regression machine learning model using data provided by the company to identify key factors that could potentially increase the lead conversion rate. This would allow the company to focus on communicating with leads who are more likely to choose a course.

# Step 1: Library Upload

I began by uploading the necessary libraries such as pandas, seaborn, matplotlib, sklearn, and statsmodels, as shown in the attached diagram.

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import statsmodels.api as sm
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn import metrics
from sklearn.metrics import precision_score, recall_score
```

✓ 6.9s

# Step 2: CSV Data Import

Next, I used the pandas library to import the CSV data into the Jupyter notebook using the `read_csv` function. I then used `leads.head()` to view the first few lines of the imported data.

```
[5] leads_csv = pd.read_csv("D:/Upgrad Class/Lead Case Study/leads.csv")
✓ 0.0s
```

[+ Code](#) [+ Markdown](#)

```
[6] leads_csv.head()
✓ 0.0s
```

..

	Prospect ID	Lead Number	Lead Origin	Lead Source	Do Not Email	Do Not Call	Converted	TotalVisits
0	7927b2df-8bba-4d29-b9a2-b6e0beafe620	660737	API	Olark Chat	No	No	0	0.0
1	2a272436-5132-4136-86fa-dcc88c88f482	660728	API	Organic Search	No	No	0	5.0
2	8cc8c611-a219-4f35-ad23-fdfd2656bd8a	660727	Landing Page Submission	Direct Traffic	No	No	1	2.0
3	0cc2df48-7cf4-4e39-9de9-19797f9b38cc	660719	Landing Page Submission	Direct Traffic	No	No	0	1.0
4	3256f628-e534-4826-9d63-4a8b88782852	660681	Landing Page Submission	Google	No	No	1	2.0

5 rows × 9 columns



# Step 3: Data Operations

I performed various operations to gain a deeper understanding of the data, such as using `leads.shape()` to determine the number of rows and columns, and `leads.dtypes()` to identify the data types of different columns. I also used `leads.info()` to view additional information about the columns.

```
leads_csv.describe()
```

	Lead Number	Converted	TotalVisits	Total Time Spent on Website	Page Views Per Visit	As A
count	9240.000000	9240.000000	9103.000000	9240.000000	9103.000000	
mean	617188.435606	0.385390	3.445238	487.698268	2.362820	
std	23405.995698	0.486714	4.854853	548.021466	2.161418	
min	579533.000000	0.000000	0.000000	0.000000	0.000000	
25%	596484.500000	0.000000	1.000000	12.000000	1.000000	
50%	615479.000000	0.000000	3.000000	248.000000	2.000000	
75%	637387.250000	1.000000	5.000000	936.000000	3.000000	
max	660737.000000	1.000000	251.000000	2272.000000	55.000000	

```
leads_csv.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9240 entries, 0 to 9239
Data columns (total 37 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Prospect ID                          9240 non-null   object
1   Lead Number                          9240 non-null   int64
2   Lead Origin                          9240 non-null   object
3   Lead Source                          9204 non-null   object
4   Do Not Email                         9240 non-null   object
5   Do Not Call                          9240 non-null   object
6   Converted                            9240 non-null   int64
7   TotalVisits                          9103 non-null   float64
8   Total Time Spent on Website          9240 non-null   int64
9   Page Views Per Visit                 9103 non-null   float64
10  Last Activity                        9137 non-null   object
11  Country                              6779 non-null   object
12  Specialization                       7802 non-null   object
13  How did you hear about X Education   7033 non-null   object
14  What is your current occupation      6550 non-null   object
15  What matters most to you in choosing a course  6531 non-null   object
16  Search                               9240 non-null   object
17  Magazine                             9240 non-null   object
18  Newspaper Article                   9240 non-null   object
```

# Step 4: Data Cleaning

I conducted data cleaning on the lead dataset. First, I identified columns with null values and calculated the percentage of null values using `is_null().sum()`. Next, I dropped columns with null values greater than 45% using the `drop` function. I then replaced null values in the remaining columns with the most frequent value. I also removed columns with unbalanced data, such as Do Not Call, Search, Magazine, etc., resulting in a balanced dataset with no null values.

✓ checking the null values in terms of percentage

+ Code

+ Markdown

```
round(100*(leads_csv.isnull().sum()/len(leads_csv.index)),2)
```

[14]

✓ 0.0s

Prospect ID	0.00
Lead Number	0.00
Lead Origin	0.00
Lead Source	0.39
Do Not Email	0.00
Do Not Call	0.00
Converted	0.00
TotalVisits	1.48
Total Time Spent on Website	0.00
Page Views Per Visit	1.48
Last Activity	1.11
Country	26.63
Specialization	36.58
How did you hear about X Education	78.46
What is your current occupation	29.11
What matters most to you in choosing a course	29.32
Search	0.00
Magazine	0.00
Newspaper Article	0.00
X Education Forums	0.00
Newspaper	0.00
Digital Advertisement	0.00
Through Recommendations	0.00
Receive More Updates About Our Courses	0.00
Tags	36.29
Lead Quality	51.59
Update me on Supply Chain Content	0.00
Get updates on DM Content	0.00
Lead Profile	74.19
City	39.71
Asymmetrique Activity Index	45.65
Asymmetrique Profile Index	45.65
Asymmetrique Activity Score	45.65
Asymmetrique Profile Score	45.65
I agree to pay the amount through cheque	0.00
A free copy of Mastering The Interview	0.00
Last Notable Activity	0.00

dtype: float64

# Step 5: Handling Outliers


I addressed outliers in numeric data columns such as TotalVisits, Total Time Spent on Website, and Page Views Per Visit to ensure data accuracy.

~ We are examining the numerical variables for any outliers.

```
leads_csv.info()
[174] ✓ 0.0s
... <class 'pandas.core.frame.DataFrame'>
RangeIndex: 9240 entries, 0 to 9239
Data columns (total 14 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Lead Origin                               9240 non-null   object
1   Lead Source                               9240 non-null   object
2   Do Not Email                             9240 non-null   object
3   Converted                                 9240 non-null   int64
4   TotalVisits                              9240 non-null   float64
5   Total Time Spent on Website              9240 non-null   int64
6   Page Views Per Visit                     9240 non-null   float64
7   Last Activity                            9240 non-null   object
8   Specialization                           9240 non-null   object
9   What is your current occupation          9240 non-null   object
10  Tags                                     9240 non-null   object
11  City                                     9240 non-null   object
12  A free copy of Mastering The Interview    9240 non-null   object
13  Last Notable Activity                    9240 non-null   object
dtypes: float64(2), int64(2), object(10)
memory usage: 1010.8+ KB
```

## Total Visits

```
plt.figure(figsize=(10,5))
sns.boxplot(y=leads_csv['TotalVisits'])
plt.show()
[175] ✓ 0.0s
...
```





# Step 6: Creating Dummy Variables

I created dummy variables for categorical columns with more than two categories. This resulted in 90 columns, preparing the data for the machine learning algorithm.

## Dummy Variable Creation

[+ Code](#) [+ Markdown](#)

```
leads_csv.info()
```

✓ 0.0s

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 9240 entries, 0 to 9239  
Data columns (total 14 columns):  
#   Column                                          Non-Null Count  Dtype    
---  ---                                            
0   Lead Origin                                   9240 non-null   object   
1   Lead Source                                   9240 non-null   object   
2   Do Not Email                                  9240 non-null   object   
3   Converted                                      9240 non-null   int64    
4   TotalVisits                                   9240 non-null   float64   
5   Total Time Spent on Website                  9240 non-null   int64    
6   Page Views Per Visit                         9240 non-null   float64   
7   Last Activity                                 9240 non-null   object   
8   Specialization                               9240 non-null   object   
9   What is your current occupation              9240 non-null   object   
10  Tags                                           9240 non-null   object   
11  City                                           9240 non-null   object   
12  A free copy of Mastering The Interview       9240 non-null   object   
13  Last Notable Activity                        9240 non-null   object   
dtypes: float64(2), int64(2), object(10)  
memory usage: 1010.8+ KB
```

```
# categorical columns  
categorical_cols= leads_csv.select_dtypes(include=['object']).columns  
categorical_cols
```

✓ 0.0s

```
Index(['Lead Origin', 'Lead Source', 'Do Not Email', 'Last Activity',  
      'Specialization', 'What is your current occupation', 'Tags', 'City',  
      'A free copy of Mastering The Interview', 'Last Notable Activity'],  
      dtype='object')
```

```
leads_csv["Do Not Email"].value_counts()
```

✓ 0.0s

```
No      8506  
Yes       734  
Name: Do Not Email, dtype: int64
```

```
leads_csv["A free copy of Mastering The Interview"].value_counts()
```

I split the data into train and test sets using the train-test library, and scaled the numeric variables using the StandardScaler method.

```
y = leads_csv['Converted']
X=leads_csv.drop('Converted', axis=1)

y.head()
```

X.head()

5 rows  $\times$  59 columns

9  0.0s Python

Python

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 6468 entries, 1871 to 5640  
Data columns (total 59 columns):  
#   Column                                     Non-Null Count  Dtype  
---  ---                                     -
```

## Step 8: Model Creation

I created models using Recursive Feature Elimination (RFE) until all the p-values of the columns were less than 0.05. I also checked for VIF values less than 3. I eliminated columns with high p-values and VIF one by one, and identified the most suitable model for further analysis. I then made predictions on the target variable and evaluated the model's accuracy, precision, recall, and specificity.

## BUILDING MODEL 1

```
X_train_sm = sm.add_constant(X_train[col])
l_model1 = sm.GLM(y_train, X_train_sm, family = sm.families.Binomial())
res = l_model1.fit()
res.summary()
```

✓ 0.0s

### Generalized Linear Model Regression Results

Dep. Variable:	Converted	No. Observations:	6468				
Model:	GLM	Df Residuals:	6452				
Model Family:	Binomial	Df Model:	15				
Link Function:	Logit	Scale:	1.0000				
Method:	IRLS	Log-Likelihood:	-1396.4				
Date:	Sun, 16 Apr 2023	Deviance:	2792.7				
Time:	01:49:15	Pearson chi2:	1.06e+04				
No. Iterations:	8	Pseudo R-squ. (CS):	0.5924				
Covariance Type:	nonrobust						
	coef	std err	z	P> z	[0.025	0.975]	
	const	-2.5631	0.088	-29.114	0.000	-2.736	-2.391
	What is your current occupation_Unemployed	2.2634	0.117	19.370	0.000	2.034	2.492
What is your current occupation_Working Professional		2.5416	0.356	7.149	0.000	1.845	3.238
	Lead Source_Welingak Website	2.9473	0.733	4.021	0.000	1.511	4.384
	Last Activity_SMS Sent	2.0669	0.109	18.950	0.000	1.853	2.281
	Tags_Already a student	-4.8662	0.718	-6.774	0.000	-6.274	-3.458
	Tags_Closed by Horizon	5.8904	1.010	5.829	0.000	3.910	7.871
	Tags_Graduation in progress	-2.5781	0.493	-5.228	0.000	-3.545	-1.612
	Tags_Interested in full time MBA	-3.8136	0.736	-5.184	0.000	-5.255	-2.372
	Tags_Interested in other courses	-3.4071	0.325	-10.479	0.000	-4.044	-2.770
	Tags_Lost to EINS	5.5233	0.727	7.601	0.000	4.099	6.948
	Tags_Not doing further education	-4.6677	1.015	-4.599	0.000	-6.657	-2.678
	Tags_Other_tags	-3.0611	0.271	-11.300	0.000	-3.592	-2.530
	Tags_Ringing	-4.4302	0.233	-19.033	0.000	-4.886	-3.974
	Tags_Will revert after reading the email	3.3450	0.183	18.256	0.000	2.986	3.704
	Tags_switched off	-4.8710	0.522	-9.334	0.000	-5.894	-3.848

```
vif = pd.DataFrame()
vif['Features'] = X_train[col].columns
vif['VIF'] = [variance_inflation_factor(X_train[col], i) for i in range(X_train[col].shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

✓ 0.1s

	Features	VIF
5	Tags_Closed by Horizon	1.33
11	Tags_Other_tags	1.29
14	Tags_switched off	1.24
10	Tags_Not doing further education	1.14
7	Tags_Interested in full time MBA	1.10
2	Lead Source_Welingak Website	1.09
6	Tags_Graduation in progress	1.09
9	Tags_Lost to EINS	1.06
1	What is your current occupation_Working Profes...	0.95
8	Tags_Interested in other courses	0.44
13	Tags_Will revert after reading the email	0.32
4	Tags_Already a student	0.29
12	Tags_Ringing	0.20
3	Last Activity_SMS Sent	0.12
0	What is your current occupation_Unemployed	0.11

# Step 9: ROC Curve

---


I used the ROC Curve to determine the best cutoff value for observations, instead of the previous 0.5 cutoff. Using this cutoff, I made predictions on the test data and calculated accuracy, precision, and specificity as the final solution for the train data. I also applied the model to the test data, which had not been seen by the model before, to assess its performance.



## Calling the ROC function

```
draw_roc(y_train_pred_final.Converted, y_train_
```


✓ 0.1s



## Plotting it

```
cutoff_df.plot.line(x='prob', y=['accuracy', 'sensi', 'speci'])  
plt.show()
```

✓ 0.1s





# Results and Conclusion

---



Based on the final machine learning model, I achieved an accuracy of 91%, sensitivity of around 81%, and specificity of around 97%. After selecting the best cutoff value of 0.3, I calculated an accuracy of 90%, sensitivity of 94%, and specificity of around 88%. Applying the model to the test data set, I concluded that the model performed well with an accuracy of 90%, sensitivity of 96%, specificity of around 87%, precision score of 83%, and recall score of 96%.



Based on these results, I recommend that the company focus on leads that are closed by 'Horizzon', lost to EINS, and avoid leads that are too